



KOTLIN USER GROUP BERLIN

One Topic to Rule 'Em All

Taming domain-event order with Kafka & Avro

Yonatan Karp-Rudin

30 min talk + 10 min Q&A

github.com/yonatankarp/one-topic-to-rule-em-all



Adventurer: Yonatan Karp-Rudin

Class: Staff Engineer · **Guild:** Billie (Berlin)

Weapons of choice: Kotlin, Spring Boot, and event-driven systems that occasionally fight back

Inventory: `ff4k` · `kotlin-design-patterns` · `kotlin-spring-boot-template`

Current quest: *making domain events arrive in the order they happened*



Yonatan

3 A.M. — YOUR PAGER

Something just leveled up a hero who doesn't exist

✦ LEVEL-UP FOR UNKNOWN ADVENTURER ad-7f3c1 —
the ledger has never heard of them!

Every event was delivered. Exactly once. In the order you produced them.

Every CRUD shop has this bug — wearing a different costume.

One topic per event type. Looks clean. Ships fast.




Type-named topics. Clear ownership. Easy diagram. Approved in review.

WHAT KAFKA ACTUALLY GUARANTEES

Ordering exists within a partition. Full stop.

partition 0



Producer writes in order → consumer reads in that exact order.

The partition is the unit of order. Nothing larger is ordered.

COROLLARY

Across topics, there is NO ordering. None.

`guild.adventurer_registered`

lag: high (busy consumer)


`guild.adventurer_leveled_up`

lag: ~zero (idle consumer)

Two topics = two independent clocks. They drift.

Star Wars shipped IV–VI before I–III. Your topics do this daily.

How the level-up overtakes the registration



THE TEMPTING NON-FIX

"Can't I just buffer and reorder by timestamp?"

- How long do you wait for the maybe-late event?
- Whose clock? Producer skew, NTP drift, GC pauses.
- What if it never comes — is it lost, or just slow?
- **You now own a stateful reordering buffer and its bugs.**

You've rebuilt a partition. Badly.

All events of one aggregate go to **one**
topic,
keyed by the

Aggregate: the thing whose events must stay in order.

Same key → same partition → guaranteed order. That's the whole trick.

BUT THE REGISTRY SAYS NO

Register a 2nd schema on guild.adventurers-value

```
POST /subjects/guild.adventurers-value/versions
{ schema: AdventurerLeveledUp }
```

409 Conflict

Schema being registered is incompatible with
an earlier schema for subject
"guild.adventurers-value"

Default subject naming assumes ONE schema per topic. A 2nd type looks like a breaking change.

A subject is not a topic

- Topic — Kafka's storage & ordering unit (the log).
- Subject — the registry's versioned compatibility scope.
- **Compatibility is checked per subject — not per topic.**
- **topic ↔ schema as 1:1 is just the DEFAULT naming strategy.**

Change the naming strategy and one topic can carry many independently-versioned schemas.

THE THREE SUBJECT-NAMING STRATEGIES

Pick the subject pattern, change everything

TopicNameStrategy

<topic>-value

guild.adventurer_registered-value

The default — and the cage.

RecordNameStrategy

<record FQN>

org.guild.adventurer.AdventurerRegistered

Schema roams any topic; niche.

TopicRecordNameStrategy

<topic>-<record FQN>

guild.adventurers-

org.guild.adventurer.AdventurerRegistered

Many schemas, scoped per topic.

WINNER

THE ONE LINE

application-fixed.yml

```
spring:
  kafka:
    producer:
      properties:
        # The one line this talk is about:
        value.subject.name.strategy:
          io.confluent.kafka.serializers.subject.TopicRecordNameStrategy
```

Producer-side only. Consumers read the schema id from the record bytes — no consumer config change.

Avro IDL — AdventurerRegistered.avdl

```
@namespace("org.guild.adventurer")
protocol AdventurerRegisteredProtocol {
    /** A new adventurer signs the guild ledger. */
    record AdventurerRegistered {
        string adventurer_id;
        string name;
        string character_class;
        timestamp_ms registered_at;
    }
}
```

Four lifecycle events. One hero.



AdventurerRegistered

the flood — 500/run



AdventurerLeveledUp

rare — the overtaker



QuestAccepted

occasional



AdventurerDied

the finale

Our hero: **Bob the Brave**. *Find the bug before my consumer does.*

MEET THE HERO

Bob the Brave

Class: Level 1 Fighter

Quest line:

Registered → LeveledUp → QuestAccepted → Died (rocks fell)

This is who we're rooting for. Watch his lifecycle stay in order — or not.





LIVE DEMO


Act 1

Two topics. Watch the level-up overtake.

→ to the terminals

WHAT YOU JUST SAW

The overtake, and the receipts



✨ LEVEL-UP FOR UNKNOWN ADVENTURER ad-7f3c1 – the ledger has never heard of them!

🙋 QuestAccepted has nowhere to go in this topology

Not a Kafka bug. Kafka did exactly what we asked: two topics, two clocks.

The class reads like the aggregate's lifecycle

```
@KafkaListener(topics = ["guild.adventurers"], groupId = "guild-ledger")
class GuildLedgerListener(/* ... */) {
    @KafkaHandler fun on(event: AdventurerRegistered) { ... }
    @KafkaHandler fun on(event: AdventurerLeveledUp) { ... }
    @KafkaHandler fun on(event: QuestAccepted) { ... }
    @KafkaHandler fun on(event: AdventurerDied) { ... }
    @KafkaHandler(isDefault = true)
    fun unknown(event: Any) { log.warn("? Unhandled: {}", ...) }
}
```



LIVE DEMO

Act 2

One topic. One line changed. Order holds.

→ to the terminals

@KafkaHandler vs. when — same hole

@KafkaHandler — framework dispatch

- One method per type.
- Spring routes by deserialized type.
- Forgotten type → silent isDefault, at runtime.

when dispatch — one explicit place

- All routing in one greppable when.
- Avro classes aren't sealed → else required.
- Forgotten type → falls to else, at runtime.

Neither gives compile-time safety on raw Avro types. For that, map events into your own sealed hierarchy at the boundary.

Both variants live in the repo — profiles
`fixed` / `fixed-when`.

Name things so the topology stays obvious

- Topics: domain.aggregate (e.g. guild.adventurers).
- Event types: fully-qualified record names (org.guild.adventurer.*).
- **One subject per (topic, type) under TopicRecordNameStrategy.**
- **Version only on a breaking change — not per release.**

The topic name is the consistency boundary. Make the boundary readable.

Let CI own the registry, not your producers

- Schemas live in a central repo, reviewed like code.
- CI registers schemas and gates compatibility per (topic, type) subject.
- **Producers run with `auto.register.schemas = false`.**
- **A breaking change fails the pipeline — not 3 a.m. in prod.**

Compatibility is a build-time gate, not a runtime surprise.

TAKEAWAYS

Four things to take home

- 1 Ordering is a partition guarantee — and nothing larger.
- 2 One aggregate → one topic, key = aggregate id.
- 3 TopicRecordNameStrategy is the knob that unlocks it.
- 4 @KafkaHandler makes the multi-type topic pleasant.

FIND ME

github.com/yonatankarp
medium.com/@yonatankarp

linkedin.com/in/yonatankarp
yonatankarp.com



github.com/yonatankarp/one-topic-to-rule-em-all



APPENDIX

Backup slides

Evolution · union schemas · consumer cost · when type-per-topic is right

Evolve one event without touching its siblings

- Compatibility is checked per subject = per (topic, event type).
- Add a field to AdventurerRegistered → only that subject evolves.
- **AdventurerLeveledUp, QuestAccepted, Died are untouched.**
- **Independent version histories per type on the same topic.**

Union schema + schema references

- Model the topic value as one union of all event types (via schema refs).
- Single subject, single schema id namespace for the topic.
- Trade-off: atomic evolution of the whole topic shape...
- **...but you lose per-type compatibility isolation.**

TopicRecordNameStrategy = many subjects, isolated. Union = one subject, atomic. Pick your coupling.

Every consumer sees every type on the topic

- All consumers deserialize all types landing on the topic.
- Filtering is consumer-side (e.g. `@KafkaHandler` / `when`).
- **Fine for aggregate streams — types are few and related.**
- **Wrong for firehose topics — many unrelated, high-volume types.**

Topic-per-event-type is correct when...

- The streams are genuinely independent — no shared aggregate.
- No ordering relationship to protect between the events.
- Heterogeneous retention or scaling per type.
- **No single aggregate id ties the records together.**

The rule isn't 'one topic always'. It's: order within an aggregate → one topic.